14/08

EX-3 : STUDY OF CLASSIFIERS WITH RESPECT TO STATISTICAL PARAMETERS

AIM:

To study of classifiers with respect to statistical parameters.

Algorithm:

=> Load an open source dataset

=> Split the dataset into training and testing sets

=> Initialize multiple classifiers

=> Train each classifiers

=> Predict on test set using each model

=> Evaluate each model using:

* Accuracy * Recall
* Precision * F1-score

=> Compare results.

PSEUDOCODE:

```
BEGIN
    load dataset
    split data into training and test
    standardize the features
    Initialize classifier
    For each classifier:
            Train
            Predict
            calculate Accuracy
            store results
    END FOR
        Display and compare
    END
```

Observation:

classifier	Decision Tree	SVM	Logistic Regression
Accuracy	84.7%	98.6%	97.4%
Avg precision	85.3%	98.7%	97.6%
Avg. Recall	83.9%	98.6%	97.3%
Avg F1-score	84.3%	98.6%	97.61%
weighted F1 score	84.7%	98.61%	97.3%

Code:

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report,
                            confusion_matrix

data = load_iris()
x = data.data
y = data.target
x_train, x_test, y_train, y_test = train_test_split
(x, y, test_size = 0.3, random_state = 43)
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
classifiers = {
    "Logistic Regression": LogisticRegression()
    "K-Neighbors": KNC(n_neighbors = 3),
    "Support Vector Machine": SVC()
    "Decision Tree": DecisionTree()
}
```

Overall:

* SVM provides best overall classification
* Decision tree shows sign of overfitting
* Logistic Regression is very close in performance.

Result:

Study of classifiers with respect to statistical parameters completed.

```
for name, clf in classifiers.items():
    clf.fit(x_train, y_train)
    ypred = clf.predict(x_test)
    print("=== {name} ===")
    print("Conf mat")
    print(confusion_matrix(y_test, y_pred))
    print(classification_report(y_test, y_pred,
      target_names: data.target_names))
```

L_3.ipynb (4) - JupyterLab

Not secure 10.1.38.19/user/ra2311047010019/lab/tree/DEEP%20LEARNING/L_3.ipynb

Verify it's you

File Edit View Run Kernel Tabs Settings Help

Logistic regression w | L_3.ipynb | L2.ipynb | EX1.ipynb | breast_cancer_bd.csv | jupyter-ra231104701

Code

Notebook Python 3 (ipykernel)

Filter files by name

/ DEEP LEARNING /

Name	Last Modified
datasets	6 days ago
EX1.ipynb	20 days ago
L_3.ipynb	next year
L2.ipynb	next year
L3.ipynb	next year
Logistic reg...	13 days ago

```
Logistic Regression
Accuracy: 0.9694444444444444
```

[25]:
```
from sklearn import metrics
```

[26]:
```
confusion_matrix=metrics.confusion_matrix(y_test,y_pred)
confusion_matrix
```

[26]:
```
array([[29,  0,  0,  0,  2,  1,  0,  0,  0,  1],
       [ 0, 22,  1,  1,  0,  0,  1,  1,  1,  1],
       [ 0,  1, 27,  2,  0,  0,  0,  1,  2,  0],
       [ 0,  0,  1, 30,  0,  0,  0,  0,  3,  0],
       [ 0,  0,  0,  1, 41,  0,  0,  4,  0,  0],
       [ 0,  0,  0,  0,  1, 43,  1,  0,  1,  1],
       [ 0,  0,  0,  0,  2,  0, 32,  0,  1,  0],
       [ 0,  0,  0,  2,  2,  0,  0, 30,  0,  0],
       [ 0,  2,  0,  2,  1,  2,  0,  0, 20,  3],
       [ 0,  0,  0,  2,  1,  0,  0,  3,  0, 34]])
```

[27]:
```
confusion_matrix=metrics.confusion_matrix(y_test,y_pred)
confusion_matrix
```

[27]:
```
array([[29,  0,  0,  0,  2,  1,  0,  0,  0,  1],
       [ 0, 22,  1,  1,  0,  0,  1,  1,  1,  1],
       [ 0,  1, 27,  2,  0,  0,  0,  1,  2,  0],
       [ 0,  0,  1, 30,  0,  0,  0,  0,  3,  0],
       [ 0,  0,  0,  1, 41,  0,  0,  4,  0,  0],
       [ 0,  0,  0,  0,  1, 43,  1,  0,  1,  1],
       [ 0,  0,  0,  0,  2,  0, 32,  0,  1,  0],
       [ 0,  0,  0,  2,  2,  0,  0, 30,  0,  0],
       [ 0,  2,  0,  2,  1,  2,  0,  0, 20,  3],
```

Simple 1 4 Python 3 (ipykernel) | Idle Mem: 283.83 MB

Mode: Command Ln 1, Col 1 L_3.ipynb 0

Search Heavier rain soon ENG 08:49

Filter files by name
/ DEEP LEARNING /

Name	Last Modified
datasets	6 days ago
EX1.ipynb	20 days ago
L_3.ipynb	next year
L2.ipynb	next year
L3.ipynb	next year
Logistic reg...	13 days ago

Logistic regression w | L_3.ipynb | L2.ipynb | EX1.ipynb | breast_cancer_bd.csv | jupyter-ra231104701
Code
Notebook
Python 3 (ipykernel)

```
[ 0,  0,  0,  2,  1,  0,  0,  3,  0, 34]])
```

[28]:
```
from sklearn.metrics import classification_report
```

[29]:
```
print(classification_report(y_test, y_pred, digits=4))
```

```
              precision    recall  f1-score   support

           0     1.0000    0.8788    0.9355        33
           1     0.8800    0.7857    0.8302        28
           2     0.9310    0.8182    0.8710        33
           3     0.7500    0.8824    0.8108        34
           4     0.8200    0.8913    0.8542        46
           5     0.9348    0.9149    0.9247        47
           6     0.9412    0.9143    0.9275        35
           7     0.7692    0.8824    0.8219        34
           8     0.7143    0.6667    0.6897        30
           9     0.8500    0.8500    0.8500        40

    accuracy                         0.8556       360
   macro avg     0.8591    0.8485    0.8515       360
weighted avg     0.8612    0.8556    0.8563       360
```

[30]:
```
print(classification_report(y_test, svm_pred, digits=4))
```

```
              precision    recall  f1-score   support

           0     1.0000    1.0000    1.0000        33
           1     1.0000    1.0000    1.0000        28
           2     1.0000    1.0000    1.0000        33
           3     1.0000    1.0000    1.0000        34
```

Logistic regression w | L_3.ipynb | L2.ipynb | EX1.ipynb | breast_cancer_bd.csv | jupyter-ra231104701
Code
Notebook Python 3 (ipykernel)
Filter files by name
/ DEEP LEARNING /
Name Last Modified
datasets 6 days ago
EX1.ipynb 20 days ago
L_3.ipynb next year
L2.ipynb next year
L3.ipynb next year
Logistic reg... 13 days ago

4	1.0000	1.0000	1.0000	46
5	0.9787	0.9787	0.9787	47
6	0.9722	1.0000	0.9859	35
7	0.9706	0.9706	0.9706	34
8	1.0000	0.9667	0.9831	30
9	0.9500	0.9500	0.9500	40
accuracy			0.9861	360
macro avg	0.9872	0.9866	0.9868	360
weighted avg	0.9862	0.9861	0.9861	360

[31]:
```
print(classification_report(y_test, lr_pred, digits=4))
```

	precision	recall	f1-score	support
0	1.0000	1.0000	1.0000	33
1	0.9655	1.0000	0.9825	28
2	0.9706	1.0000	0.9851	33
3	0.9706	0.9706	0.9706	34
4	1.0000	0.9565	0.9778	46
5	0.9167	0.9362	0.9263	47
6	0.9444	0.9714	0.9577	35
7	1.0000	0.9706	0.9851	34
8	0.9667	0.9667	0.9667	30
9	0.9744	0.9500	0.9620	40
accuracy			0.9694	360
macro avg	0.9709	0.9722	0.9714	360
weighted avg	0.9699	0.9694	0.9695	360

L_3.ipynb (4) - JupyterLab

Not secure 10.1.38.19/user/ra2311047010019/lab/tree/DEEP%20LEARNING/L_3.ipynb

Verify it's you

File Edit View Run Kernel Tabs Settings Help

Logistic regression w | L_3.ipynb | L2.ipynb | EX1.ipynb | breast_cancer_bd.csv | jupyter-ra231104701

Code

Notebook Python 3 (ipykernel)

Filter files by name

/ DEEP LEARNING /

Name	Last Modified
datasets	6 days ago
EX1.ipynb	20 days ago
L_3.ipynb	next year
L2.ipynb	next year
L3.ipynb	next year
Logistic reg...	13 days ago

[1]:
```
from sklearn.datasets import load_digits
```

[2]:
```
d=load_digits()
```

[3]:
```
x=d.data
y=d.target
```

[4]:
```
x
```

[4]:
```
array([[ 0.,  0.,  5., ...,  0.,  0.,  0.],
       [ 0.,  0.,  0., ..., 10.,  0.,  0.],
       [ 0.,  0.,  0., ..., 16.,  9.,  0.],
       ...,
       [ 0.,  0.,  1., ...,  6.,  0.,  0.],
       [ 0.,  0.,  2., ..., 12.,  0.,  0.],
       [ 0.,  0., 10., ..., 12.,  1.,  0.]])
```

[5]:
```
y
```

[5]:
```
array([0, 1, 2, ..., 8, 9, 8])
```

[6]:
```
from sklearn.model_selection import train_test_split
```

[7]:
```
x_train,x_test,y_train,y_test=train_test_split(x,y, test_size=0.2,random_state=42)
```

[8]:
```
from sklearn.tree import DecisionTreeClassifier
```

[9]:
```
clf=DecisionTreeClassifier()
```

Simple 1 4 Python 3 (ipykernel) | Idle Mem: 283.83 MB

Mode: Command Ln 1, Col 1 L_3.ipynb 0

Search

Gold +0.99%

ENG 08:49